


Regularization



METIS

OLS Regression Results

Dep. Variable:	DomesticTotalGross	R-squared:	0.286
Model:	OLS	Adj. R-squared:	0.278
Method:	Least Squares	F-statistic:	34.82
Date:	Sun, 14 Sep 2014	Prob (F-statistic):	6.80e-08
Time:	21:59:46	Log-Likelihood:	-1738.1
No. Observations:	89	AIC:	3480.
Df Residuals:	87	BIC:	3485.
Df Model:	1		

	coef	std err	t	P> t 	[95.0% Conf. Int.]
Budget	0.7846	0.133	5.901	0.000	0.520 1.049
Ones	4.44e+07	1.27e+07	3.504	0.001	1.92e+07 6.96e+07

Omnibus:	39.749	Durbin-Watson:	0.674
Prob(Omnibus):	0.000	Jarque-Bera (JB):	99.441
Skew:	1.587	Prob(JB):	2.55e-22
Kurtosis:	7.091	Cond. No.	1.54e+08

$$AIC = 2k - 2\ln(L)$$



#

parameters


Log

likelihood

While awarding goodness of fit, penalize model complexity

While awarding goodness of fit, penalize model complexity

Why not do that while we are fitting?

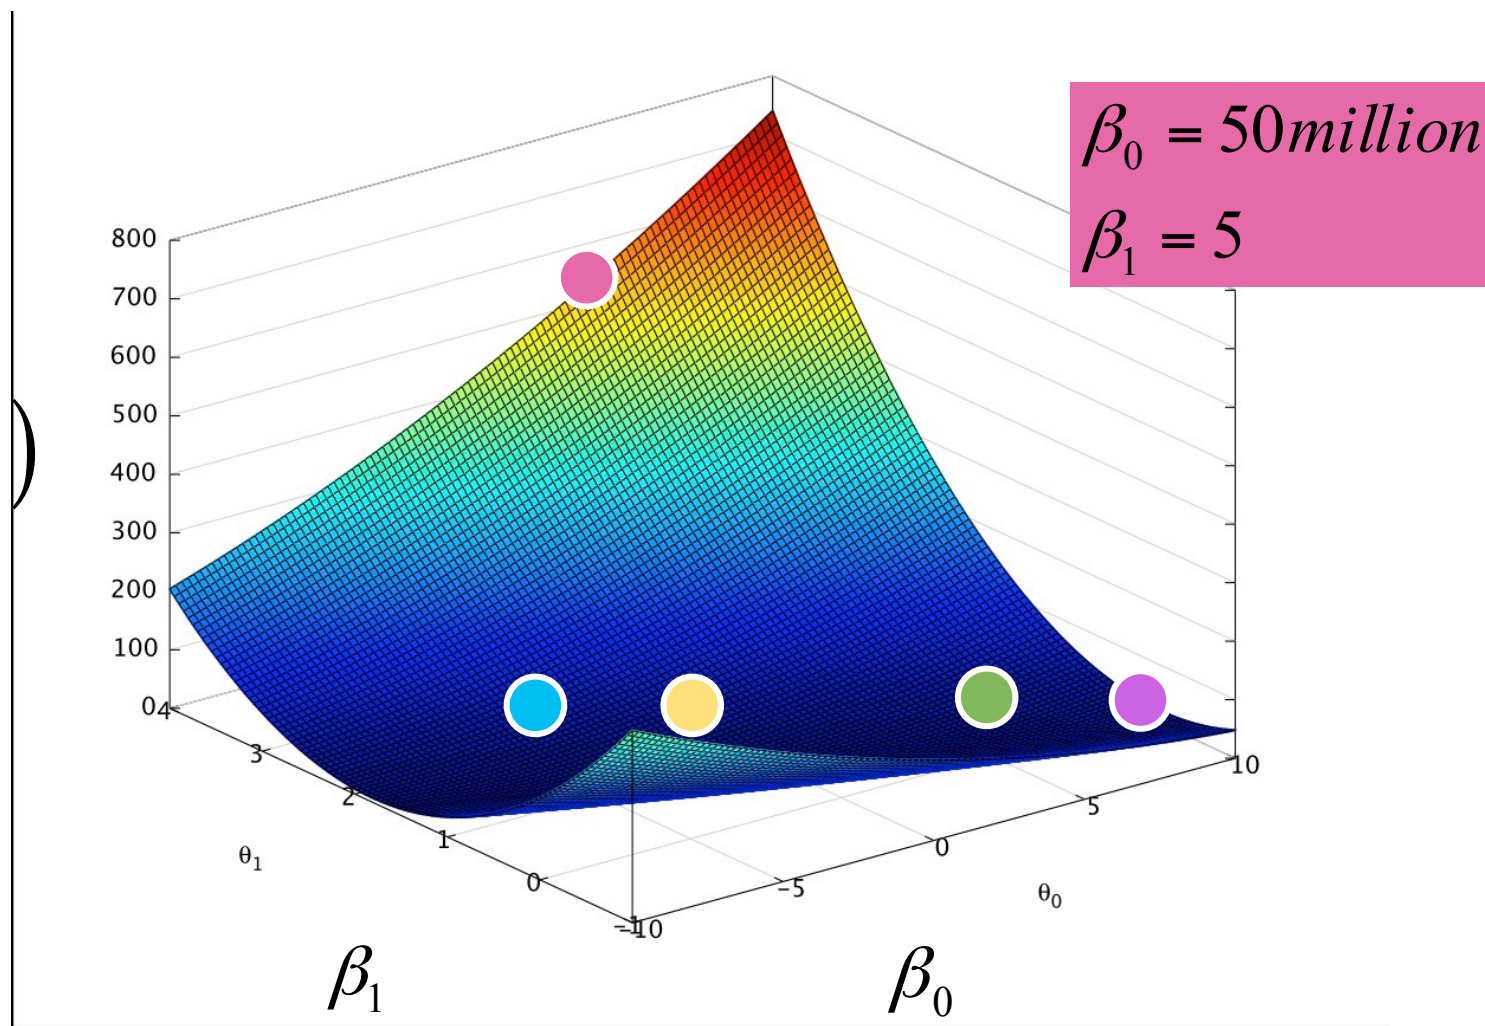


Cost function

Takes a model (specific parameter values), returns a score

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left((\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

$$J(\beta_0, \beta_1)$$



$$\beta_0 = 80\text{million}$$

$$\beta_1 = 0.5$$

$$\beta_0 = 0$$

$$\beta_1 = 1.5$$

$$\beta_0 = 120\text{million}$$

$$\beta_1 = 0.1$$

$$\beta_0 = 30\text{million}$$

$$\beta_1 = 2$$

Cost function

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left((\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



Lower for
better fits


Cost function

Add a penalty for the size of each parameter!

$$J(\beta_0, \beta_1) = \underbrace{\frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2}_{\substack{\text{Low: good fit} \\ \text{High: bad fit}}} + \underbrace{\lambda \sum_{j=1}^k \beta_j^2}_{\substack{\text{Low: simple model} \\ \text{High: complex model}}}$$

Diagnostics to detect under/overfitting


Underfitting



$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m (y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)})^2 + \lambda \sum_{j=1}^k \beta_j^2$$

J = V. High + Low


Just Right



$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m (y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)})^2 + \lambda \sum_{j=1}^k \beta_j^2$$

J = Low + Low

Overfitting



$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m (y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)})^2 + \lambda \sum_{j=1}^k \beta_j^2$$

J = Low + V. High

Ridge Regression

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

Underfitting

$$J = \text{V. High} + \text{Low}$$


Just Right

$$J = \text{Low} + \text{Low}$$

Overfitting

$$J = \text{Low} + \text{V. High}$$

$$\lambda = 1$$


$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

$$y_{\beta}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \varepsilon$$

Underfitting

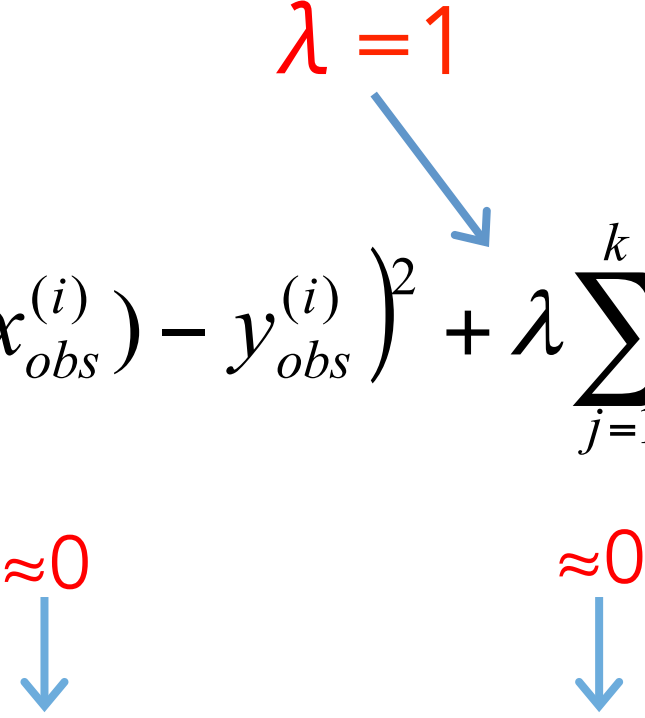
$$J = \text{V. High} + \text{Low}$$

Just Right

$$J = \text{Low} + \text{Low}$$

Overfitting

$$J = \text{Low} + \text{V. High}$$

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$


$$y_{\beta}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \varepsilon$$

Underfitting

$J = \text{V. High} + \text{Medium}$


Just Right

$J = \text{Low} + \text{V High}$

Overfitting

$J = \text{Low} + \text{V V High}$

VERY LARGE
underfit


$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

≈ 0



≈ 0



≈ 0



≈ 0



$$y_{\beta}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \varepsilon$$

Underfitting

$$J = \text{V. High} + \text{Tiny}$$

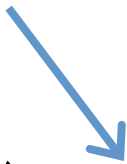
Just Right

$$J = \text{Low} + \text{Tiny}$$

Overfitting


$$J = \text{Low} + \text{Tiny}$$


very small
possible
overfit


$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

$$y_{\beta}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \varepsilon$$

Error vs. regularization λ





Ridge Regularization (L2)

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

Ridge Regularization (L2)

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

Lasso Regularization (L1)

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k |\beta_j|$$

Ridge Regularization (L2)

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k \beta_j^2$$

Lasso Regularization (L1)

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda \sum_{j=1}^k |\beta_j|$$

Elastic Net (L1 + L2)

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left(y_{\beta}(x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 + \lambda_1 \sum_{j=1}^k |\beta_j| + \lambda_2 \sum_{j=1}^k \beta_j^2$$

How can I use this?

We were doing:

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X,Y)
```

How can I use this?

We were doing:

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X,Y)
```

To use Lasso Regularization:

```
from sklearn.linear_model import Lasso  
model = Lasso(1.0)  
model.fit(X,Y)
```

λ (sklearn Calls It
alpha)



How can I use this?

We were doing:

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X,Y)
```

To use Ridge Regularization:

```
from sklearn.linear_model import Ridge  
model = Ridge(1.0)  
model.fit(X,Y)
```

λ (sklearn Calls It alpha)



How can I use this?

We were doing:

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X,Y)
```

To use Elastic Net:

```
from sklearn.linear_model import ElasticNet  
model = ElasticNet(1.0, l1_ratio = 0.5)  
model.fit(X,Y)
```



total weight for the full
penalty term



ratio of l1/l2 penalty

How can I use this?

We were doing:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit(Y, X)
```


How can I use this?

We were doing:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit(Y, X)
```

To use Lasso Regularization:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit_regularized(Y, X, alpha=1.0)
```

How can I use this?

We were doing:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit(Y, X)
```

To use Ridge Regularization:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit_regularized(Y, X,  
                                alpha=1.0,  
                                L1_wt = 0)
```

How can I use this?

We were doing:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit(Y, X)
```

To use Elastic Net:

```
import statsmodels.formula.api as sm  
model = sm.OLS.fit_regularized(Y, X,  
                                alpha=1.0,  
                                L1_wt = 0.5)
```

My model is not
awesome
enough.

What do I do?

Try these and check test error again:

Use a smaller set of features

Try adding polynomials

Check functional forms for each feature

Try including other features

Use more data (bigger training set)

Regularization

Try these and check test error again:

Use a smaller set of features

Try adding polynomials

Check functional forms for each feature

Try including other features

Use more data (bigger training set)

Regularization: Increase/decrease λ